

Probabilistic Programming Languages

Build Your Owl PPL

Guillaume Baudart

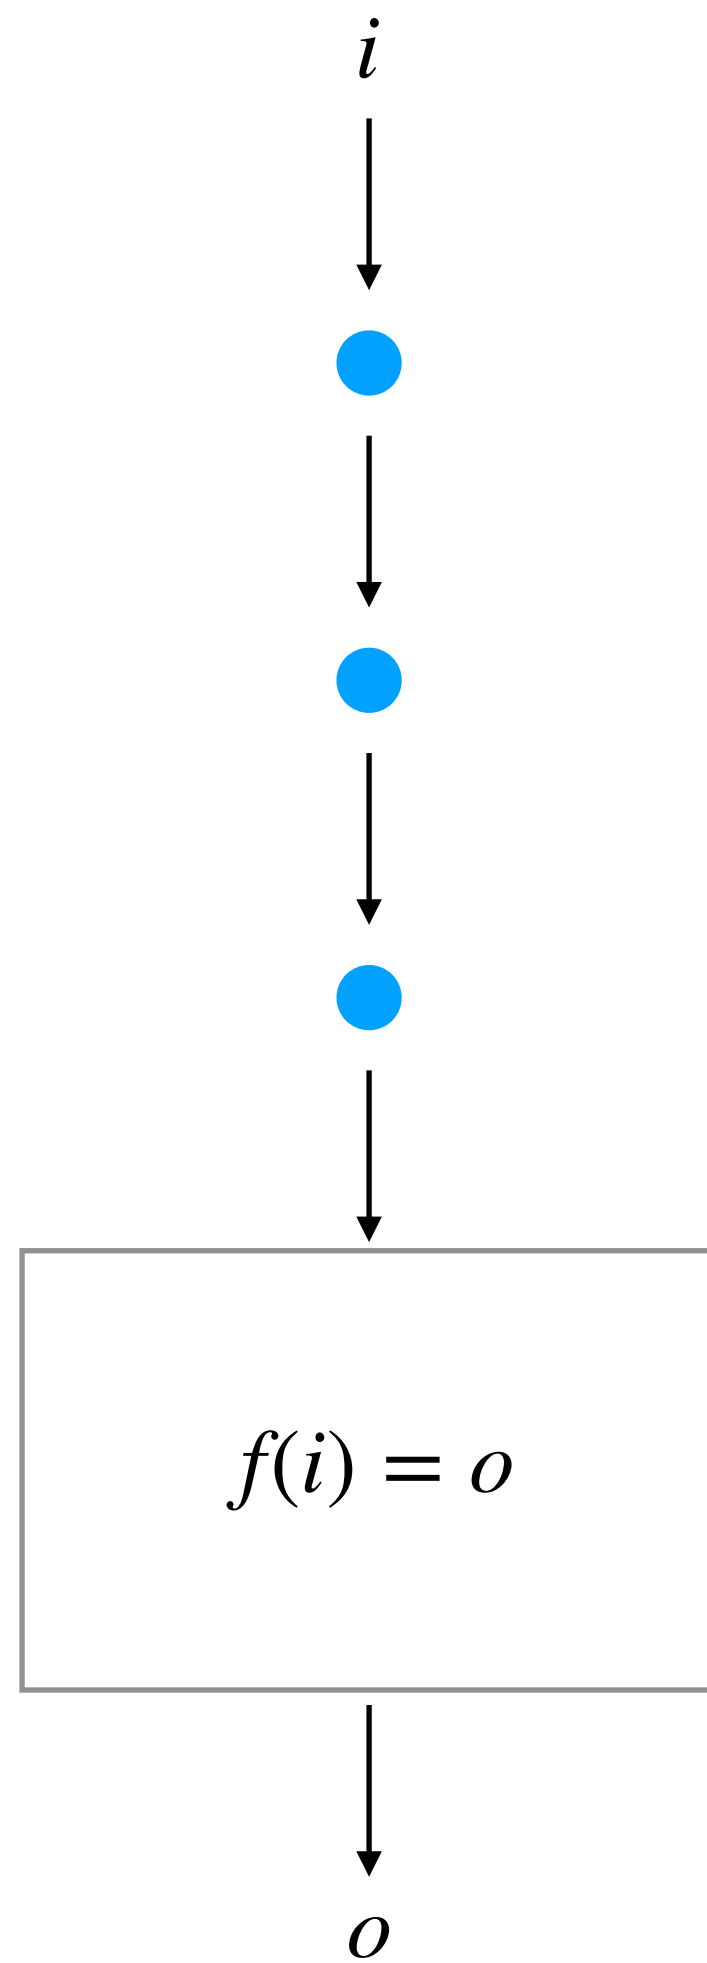
MPRI 2023-2024

Reminders


BYO-PPL

infer : $(\alpha \rightarrow \beta) \rightarrow \alpha \rightarrow \beta$ dist

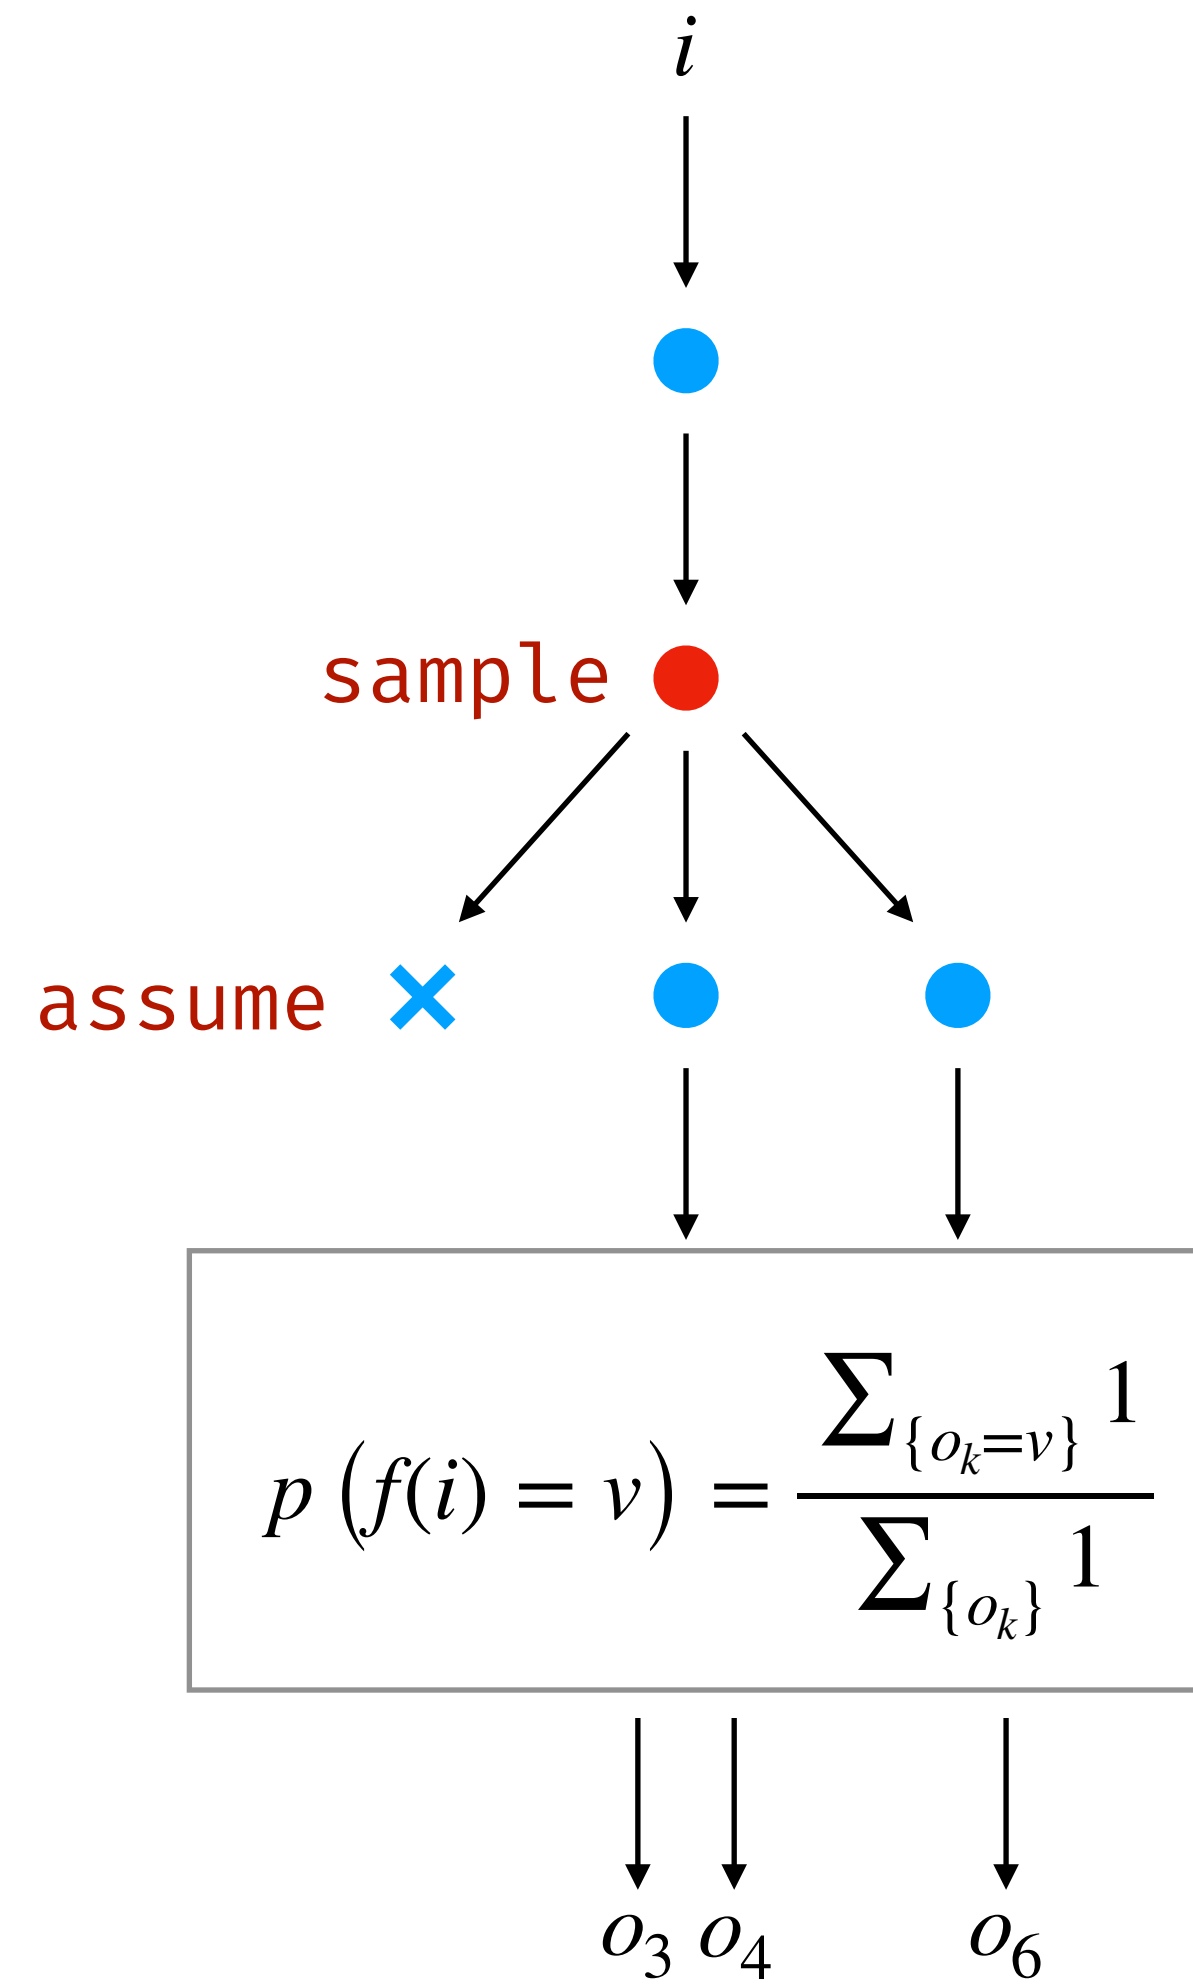
program



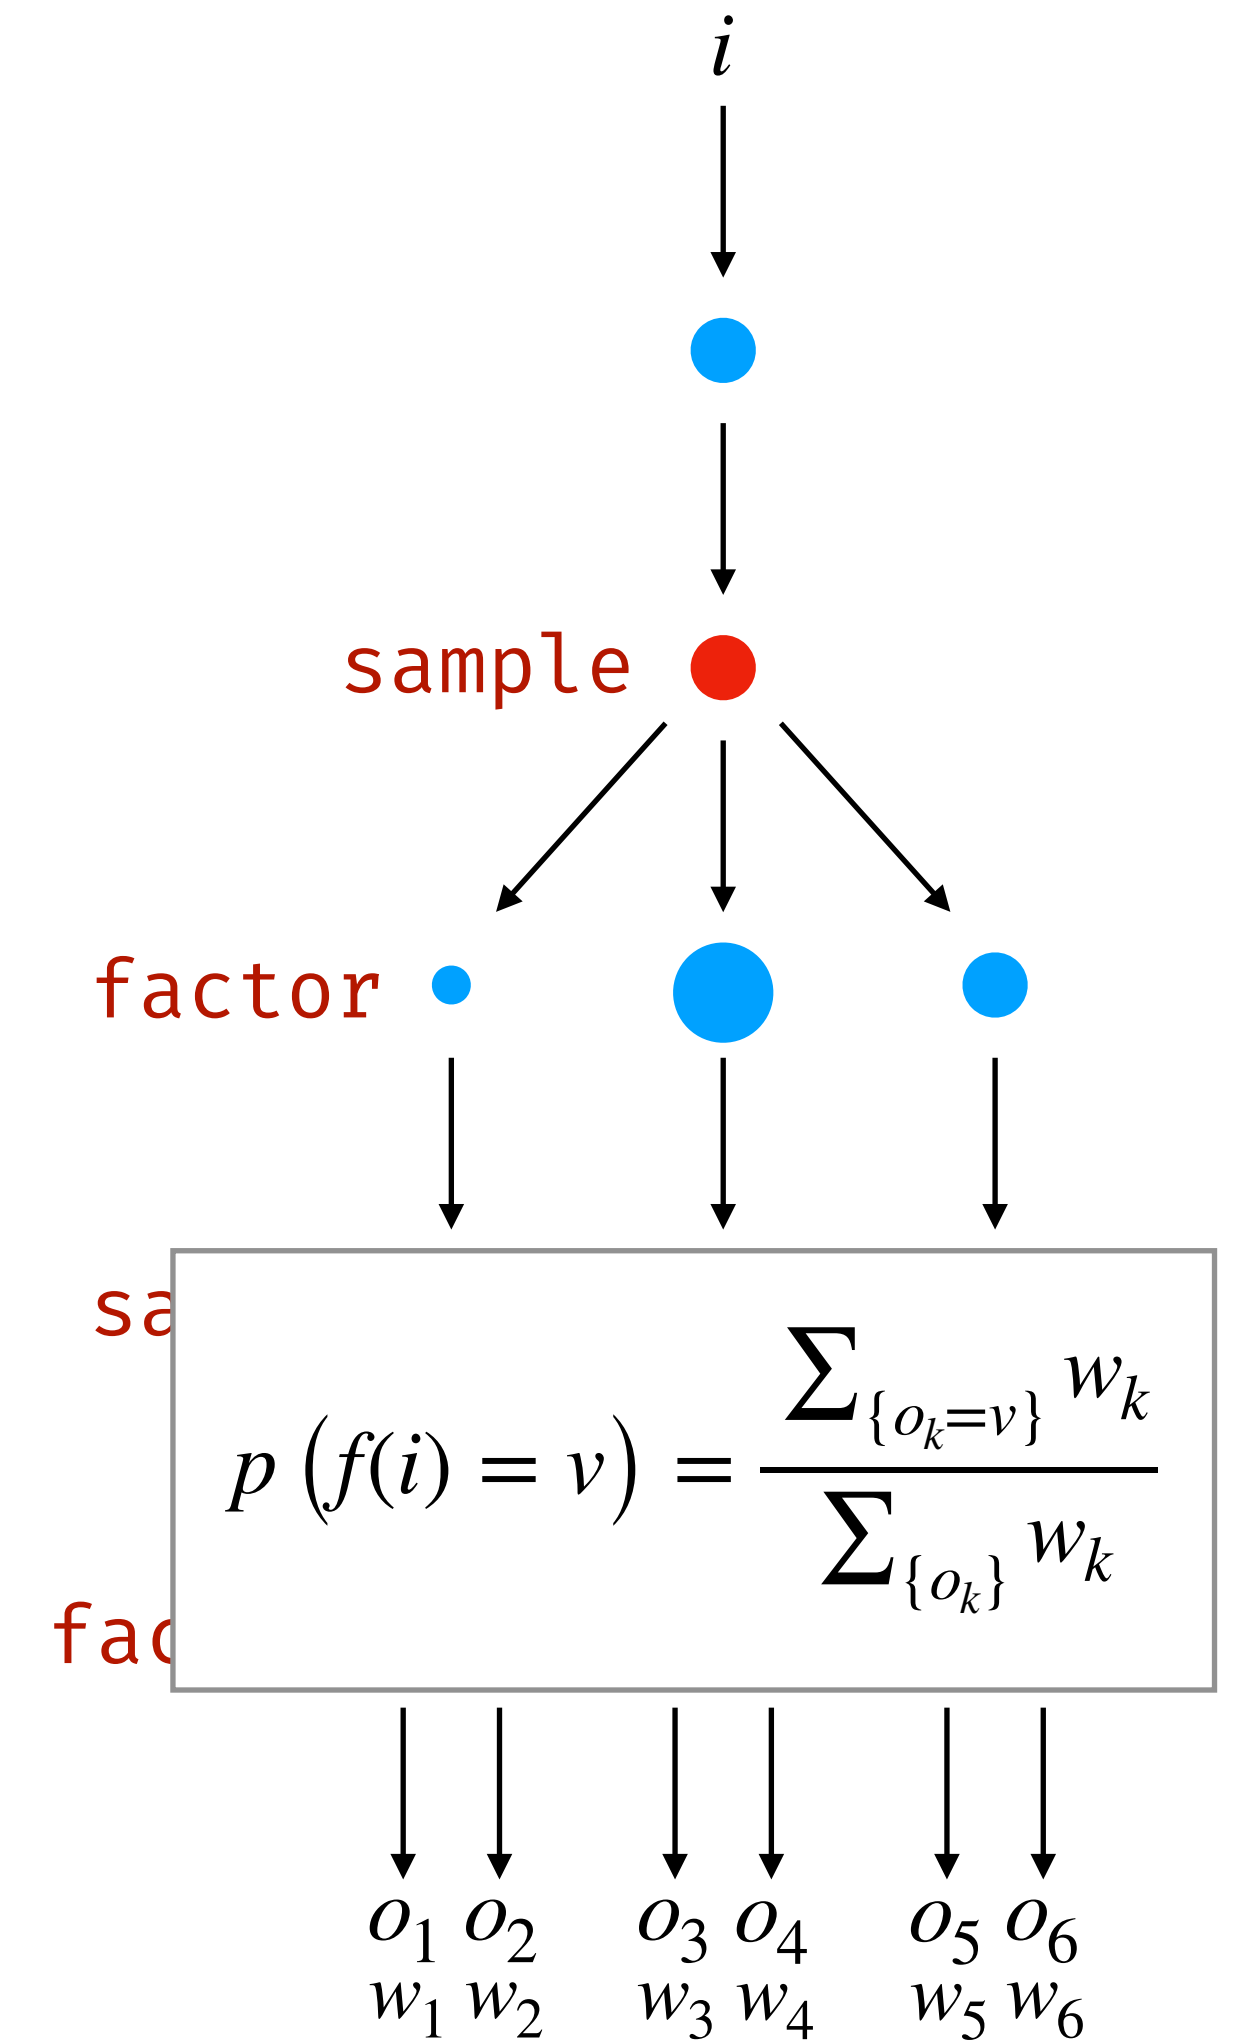
sample




assume



factor



Semantics: (un)normalized measures



$$\begin{aligned}
 \{\!\{e}\!\}_\gamma^\phi &= \lambda U. \delta_{\llbracket e \rrbracket_\gamma^\phi}(U) \text{ if } \text{kindOf}(e) = D \\
 \{\!\{\text{let } p = e_1 \text{ in } e_2\}\!\}_\gamma^\phi &= \lambda U. \int_{\llbracket \text{typeOf}(e_1) \rrbracket} \{\!\{e_1\}\!\}_\gamma^\phi(dv) \{\!\{e_2\}\!\}_{\gamma+[p \leftarrow v]}^\phi \\
 \{\!\{\text{sample}(e)\}\!\}_\gamma^\phi &= \lambda U. \llbracket e \rrbracket_\gamma^\phi(U) \\
 \{\!\{\text{factor}(e)\}\!\}_\gamma^\phi &= \lambda U. \llbracket e \rrbracket_\gamma^\phi \times \delta_{(\cdot)}(U) \\
 \{\!\{\text{observe}(e_1, e_2)\}\!\}_\gamma^\phi &= \lambda U. \text{pdf}(\llbracket e_1 \rrbracket_\gamma^\phi)(\llbracket e_2 \rrbracket_\gamma^\phi) \cdot \delta_{(\cdot)}(U) \\
 \llbracket \text{infer}(e) \rrbracket_\gamma^\phi &= \begin{cases} \frac{\lambda U. \{\!\{e\}\!\}_\gamma^\phi(U)}{\{\!\{e\}\!\}_\gamma^\phi(\llbracket \text{typeOf}(e) \rrbracket)} \\ \text{Error if } \{\!\{e\}\!\}_\gamma^\phi(\llbracket \text{typeOf}(e) \rrbracket) \in \{0, \infty\} \end{cases}
 \end{aligned}$$

Exercises

Prove the following properties

■ `sample mu (* where mu is defined on [a, b] *)`

`≡`

`let x = sample (uniform (a, b)) in`

`let () = observe (mu, x) in`

`x`

■ `observe (mu, x) (* where mu is a discrete distribution *)`

`≡`

`let y = sample mu in`

`assume x = y`

■ `sample (bernoulli (0.5))`

`≡`

`let x = sample (gaussian (0., 1.)) in`

`x > 0`

Build Your Owl PPL

BYO-PPL

Outline

For a given inference algorithm, how to implement `sample`, `assume`, `factor`, `observe`, and `infer`?

I - Problem with basic inference

- Curse of dimensionality
- Resampling and checkpoints

II - Continuation Passing Style (CPS) models

III - Revisiting inference with CPS programming

- Sample generation
- Importance sampling
- Particle filter

IV - Inference formalization

- Weighted samplers
- Big-step semantics with checkpoints

HMM: Hidden Markov Model

Track the position of an agent from noisy observations

- The current position should not be too far from the previous position
- The observations should not be too far from the current position

Probabilistic model: $\forall t \in \mathbb{N}$

- $x_t \sim \mathcal{N}(x_{t-1}, \text{speed})$
- $y_t \sim \mathcal{N}(x_t, \text{noise})$

HMM: Hidden Markov Model

hmm.ml

```
open Basic.Importance_sampling

let hmm prob data =
  let rec gen states data =
    match (states, data) with
    | [], y :: data → gen [ y ] data
    | states, [] → states
    | pre_x :: _, y :: data →
      let x = sample prob (gaussian ~mu:pre_x ~sigma:1.0) in
      let () = observe prob (gaussian ~mu:x ~sigma:1.0) y in
      gen (x :: states) data
  in
  gen [] data


let _ =
  let data = Owl.Arr.linspace 0. 20. 20 ▷ Owl.Arr.to_array ▷ Array.to_list in
  let dist = Distribution.split_list (infer hmm data) in
  let m_x = List.rev (List.map Distribution.mean dist) in
  List.iter2 (Format.printf "%f >> %f@.") data m_x
```

HMM: Hidden Markov Model

```
› dune exec ./examples/hmm.exe
```

```
0.000000 >> 0.000000
1.052632 >> 0.278989
2.105263 >> 2.923428
3.157895 >> 2.812035
4.210526 >> 2.328341
5.263158 >> 1.742109
6.315789 >> 2.518105
7.368421 >> 3.958375
8.421053 >> 5.946233
9.473684 >> 7.329554
10.526316 >> 9.293653
11.578947 >> 10.181831
12.631579 >> 8.549409
13.684211 >> 9.323073
14.736842 >> 9.280692
15.789474 >> 9.352218
```

...



HMM: Importance sampling

Problem:


- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance sampling

Problem:


- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance sampling

Problem:


- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance sampling

Problem:


- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



HMM: Importance sampling

Problem:

- Each particle does a random walk and compute a score
- No re-calibration based on observations
- Score decreases at each factor statement



The curse of dimensionality

Problem becomes harder as the dimension increases

Basic inference: rejection sampling, importance sampling

- Performances decrease exponentially when the dimension increases
- Only use for low-dimension models

17h45mn

How to mitigate this problem?

- Make assumptions about the posterior distributions
- Break the problem into simpler, smaller problems



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step


- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step

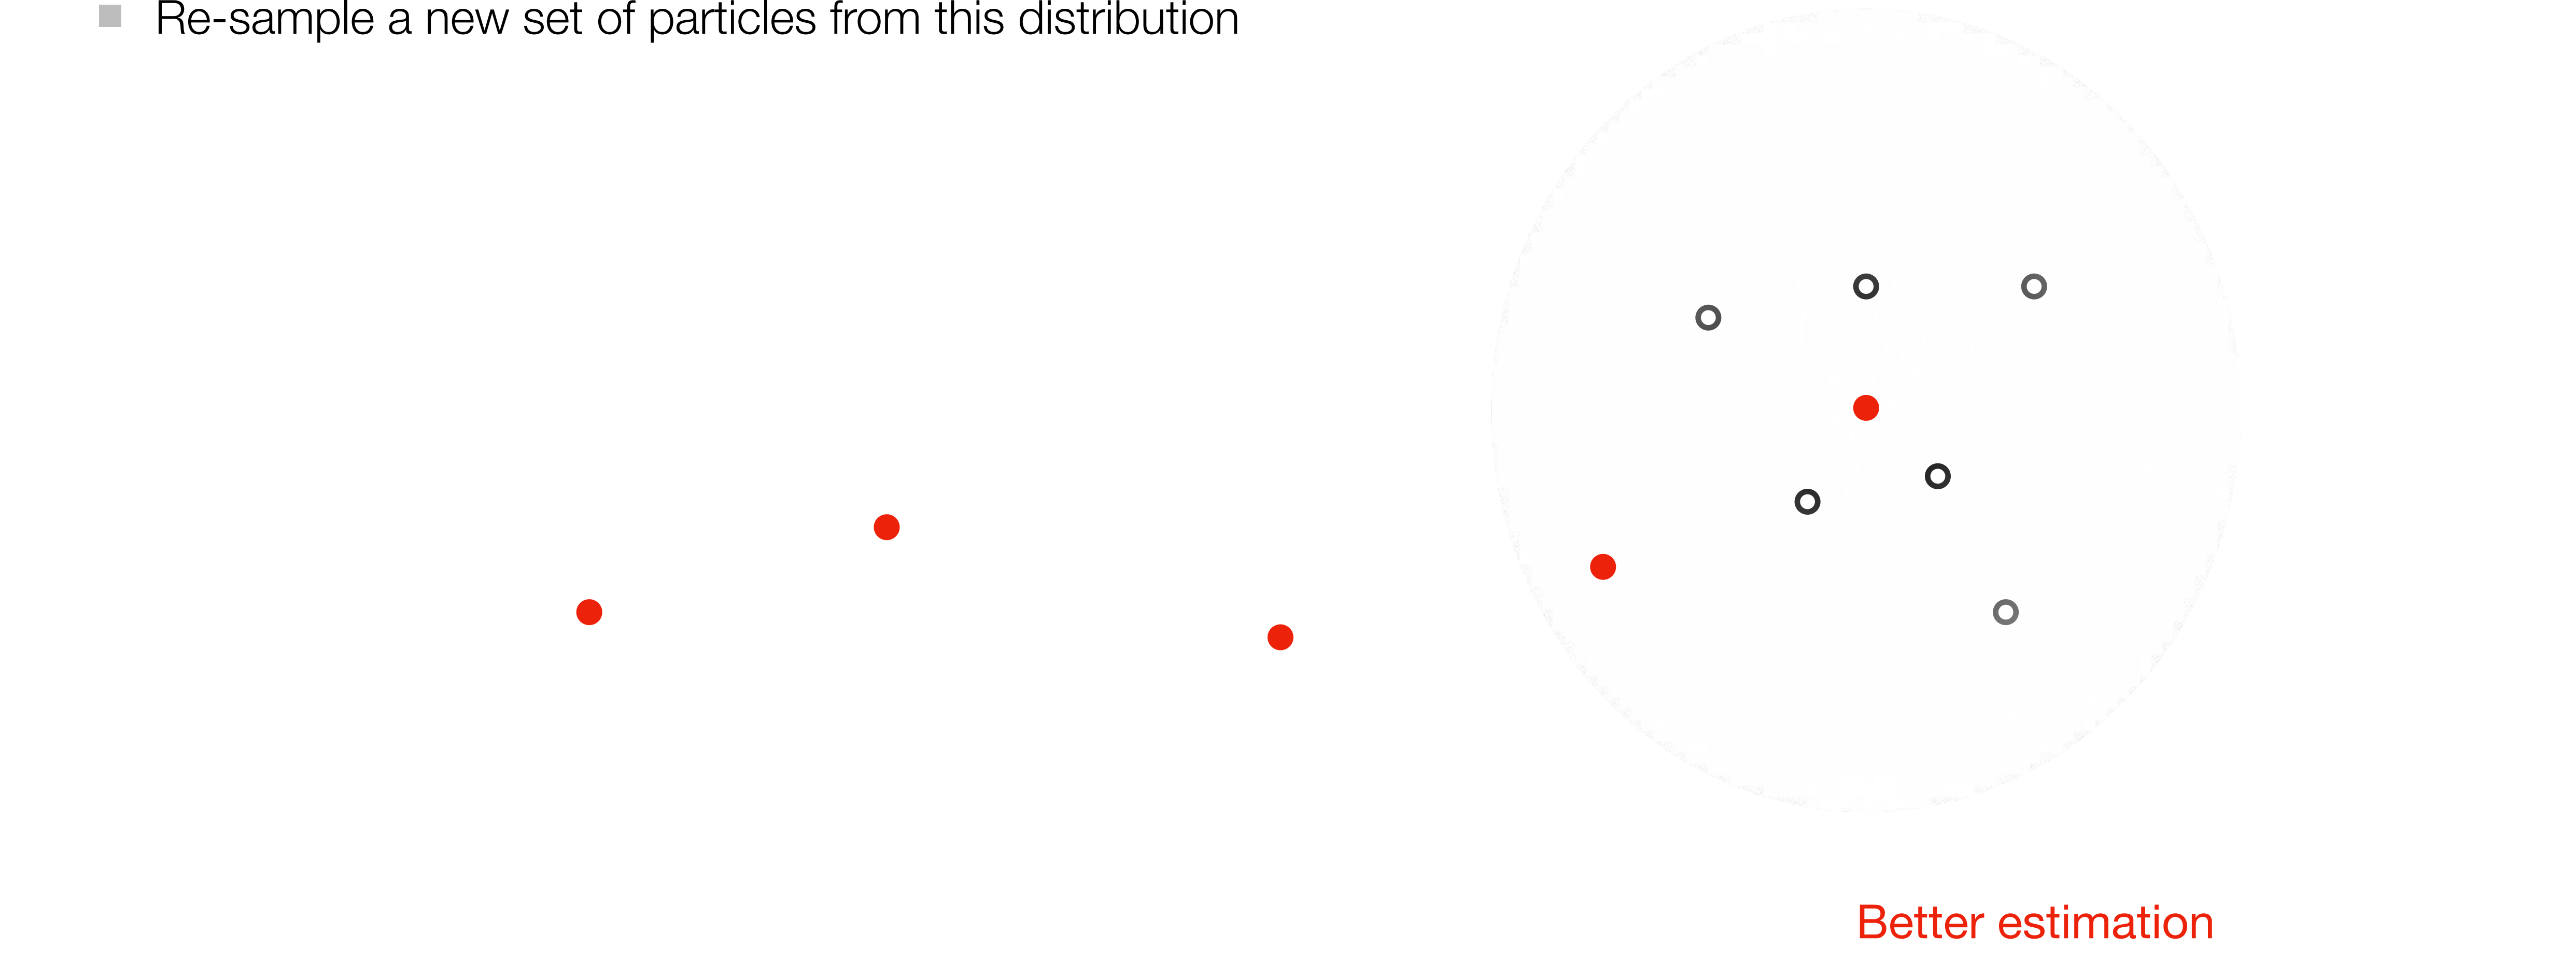
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



HMM: Particle filter

Add a resampling step

- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution



Problem: Duplications

How can we duplicate a particle during execution?

- Rerun the particle from the start?
- Force reuse sampled values?
- Clone the memory state?

Continuation Passing Style

- Functions take an extra argument `k`: the continuation
- `k` implements what should be done with the result of the function
- In our context, we can use continuation to interrupt/restart the execution of a model

Continuation Passing Style (CPS)

BYO-PPL

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) →  
    let hl = tree_height l in  
    let hr = tree_height r in  
    (1 + max hl hr)
```

1. Add intermediate values

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t k =  
  match t with  
  | Empty → k 0  
  | Node (_, l, r) →  
    let hl = tree_height l in  
    let hr = tree_height r in  
    k (1 + max hl hr)
```

1. Add intermediate values
2. Add call to continuation

Reminders: CPS

tree.ml

```
let rec tree_height t =  
  match t with  
  | Empty → 0  
  | Node (_, l, r) → 1 + max (tree_height l) (tree_height r)
```

```
let rec tree_height t k =  
  match t with  
  | Empty → k 0  
  | Node (_, l, r) →  
    tree_height l (fun hl →  
      tree_height r (fun hr →  
        k (1 + max hl hr)))
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

Funny bernoulli CPS

funny_bernoulli.ml

```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

Funny bernoulli CPS

funny_bernoulli.ml

```
let funny_bernoulli () =  
  let a = sample (bernoulli ~p:0.5) in  
  let b = sample (bernoulli ~p:0.5) in  
  let c = sample (bernoulli ~p:0.5) in  
  let () = assume (a = 1 || b = 1) in  
  a + b + c
```

```
let funny_bernoulli () k =  
  sample (bernoulli ~p:0.5) (fun a →  
    sample (bernoulli ~p:0.5) (fun b →  
      sample (bernoulli ~p:0.5) (fun c →  
        assume (a = 1 || b = 1) (fun () →  
          k (a + b + c))))
```

1. Add intermediate values
2. Add call to continuation
3. Turn let/in into nested function call

CPS monadic operators

cps_operators.ml

```
let return e k = k e
```

```
let ( let* ) e f k = e (fun x → f x k)    (* let* x = e in f(x) *)
```

CPS monadic operators

cps_operators.ml

```
let return e k = k e
```

```
let ( let* ) e f k = e (fun x → f x k)    (* let* x = e in f(x) *)
```

```
let funny_bernoulli () k =  
  sample (bernoulli ~p:0.5) (fun a →  
    sample (bernoulli ~p:0.5) (fun b →  
      sample (bernoulli ~p:0.5) (fun c →  
        assume (a = 1 || b = 1) (fun () →  
          k (a + b + c))))
```

```
let funny_bernoulli () =  
  let* a = sample (bernoulli ~p:0.5) in  
  let* b = sample (bernoulli ~p:0.5) in  
  let* c = sample (bernoulli ~p:0.5) in  
  let* () = assume (a = 1 || b = 1) in  
  return (a + b + c)
```

Sample generation (CPS)

BYO-PPL

CPS models

infer.ml

```
module Gen : sig
  type 'a prob
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  val sample : 'a Distribution.t → ('a → 'b next) → 'b next
  val factor : float → (unit → 'b next) → 'b next
  val draw: ('a, 'b) model → 'a → 'b
end = struct ... end
```


Type 'a prob

- Store all information required for inference (e.g., particles array)
- Type ('a, 'b) model capture input/output types

Models where probabilistic constructs are CPS functions

- Two arguments: input 'a and a continuation on the return value ('b → 'b next).
- The return value is a continuation 'a next that updates a probabilistic state of type 'a prob.

Sample generation



Sample generation

infer.ml

```
module Gen = struct
  type 'a prob = 'a option
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  let exit v _prob = Some v

  let sample d k prob =
    let v = Distribution.draw d in
    k v prob

  let factor _s k prob = k () prob

  let draw m data =
    let v = (m data) exit None in
    Option.get v
end
```

Funny bernoulli

funny_bernoulli.ml

```
open Infer.Gen
```

```
let funny_bernoulli () =  
  let* a = sample (bernoulli ~p:0.5) in  
  let* b = sample (bernoulli ~p:0.5) in  
  let* c = sample (bernoulli ~p:0.5) in  
  let* () = assume (a = 1 || b = 1) in  
  return (a + b + c)
```

```
let _ =  
  for _ = 1 to 10 do  
    let v = draw funny_bernoulli () in  
    Format.printf "%d " v  
  done
```

```
> dune exec ./examples/funny_bernoulli.exe
```

```
1 1 2 2 2 2 2 1 3 2
```

Importance sampling (CPS)

BYO-PPL

Importance sampling

infer.ml

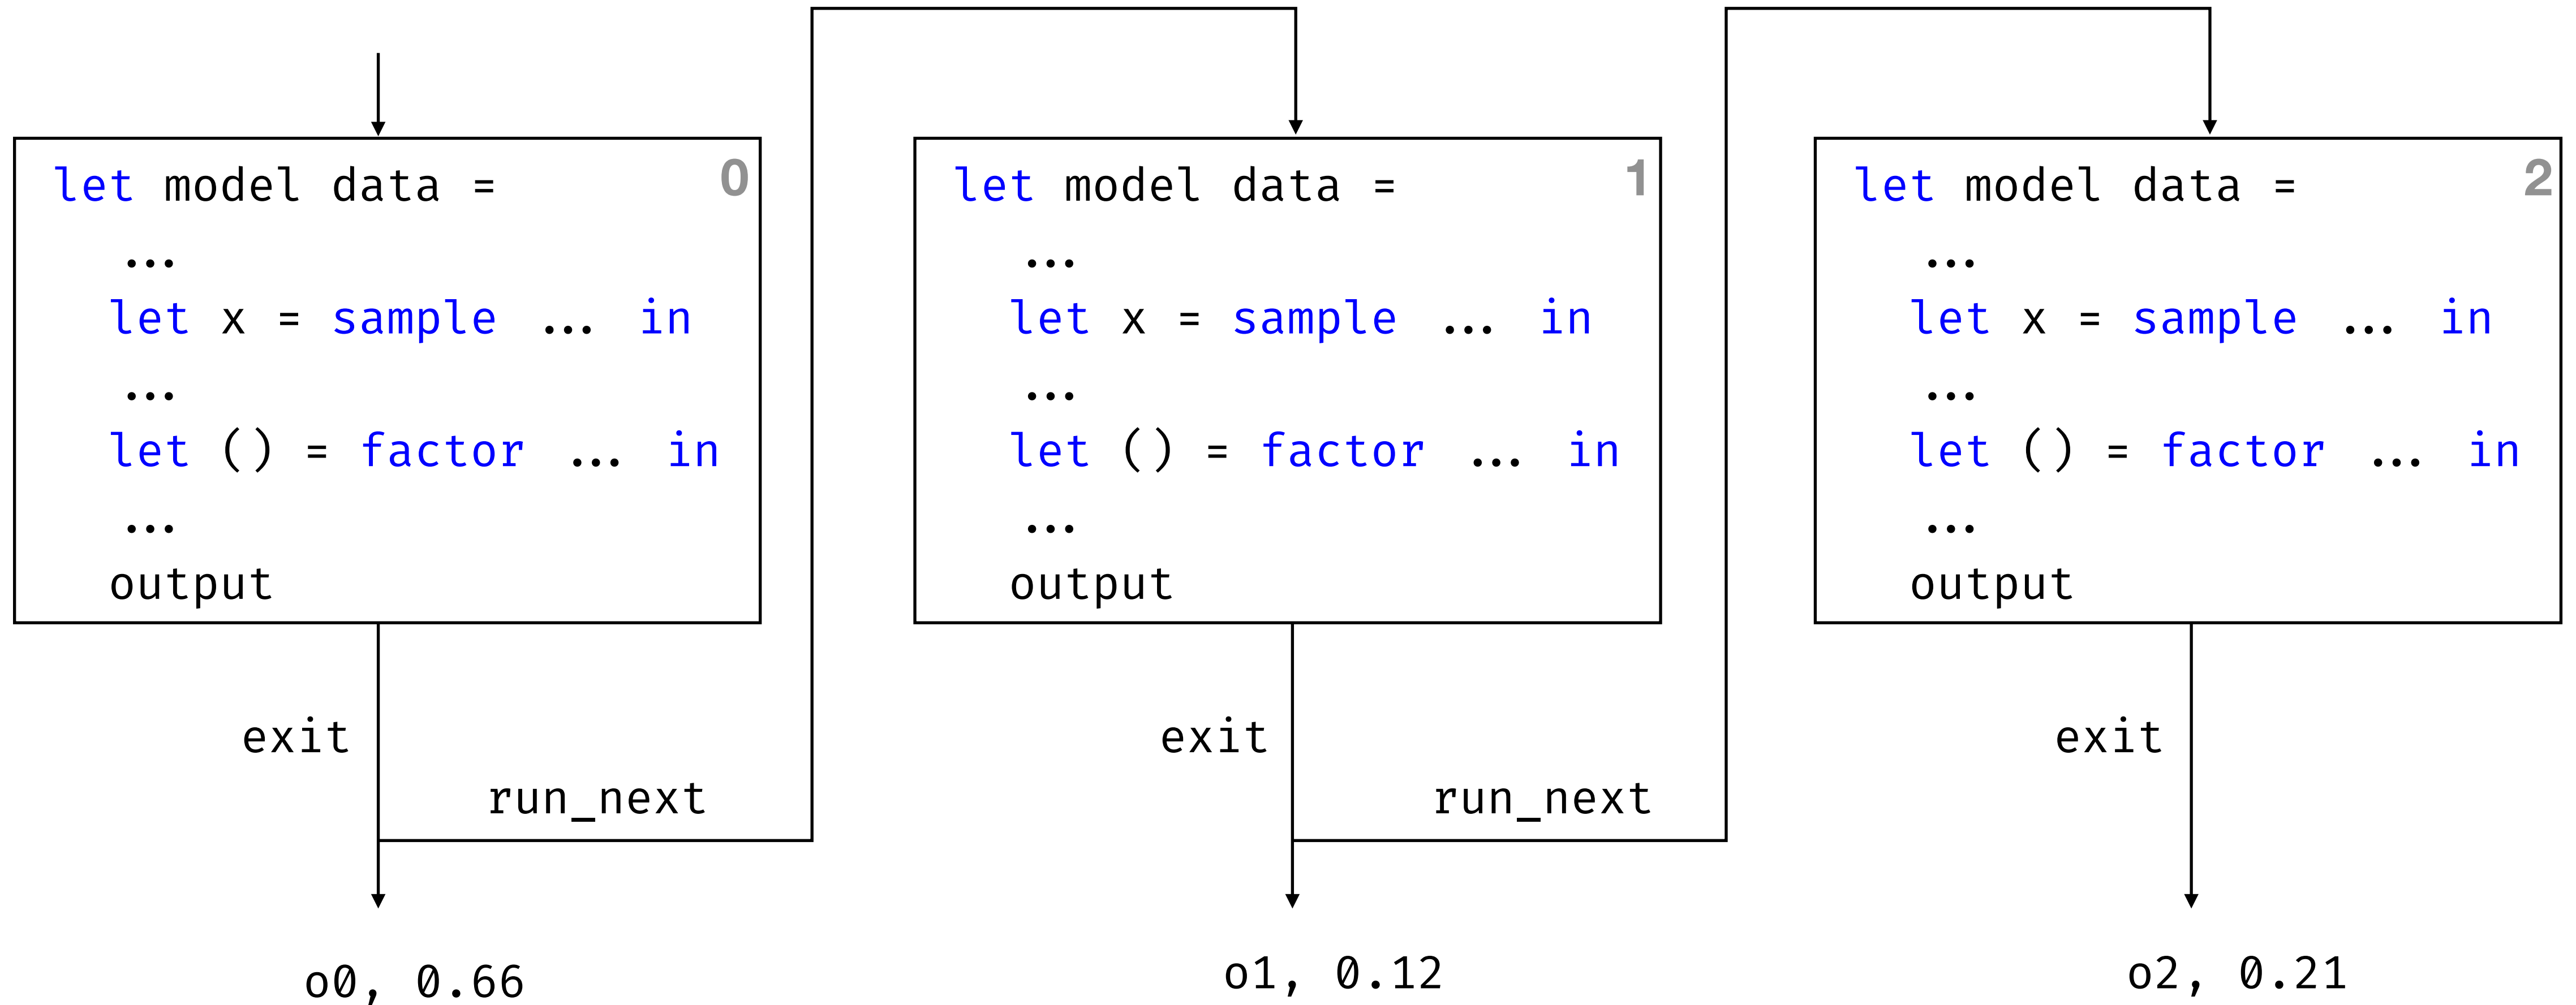
```
module Importance_sampling : sig
  type 'a prob
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  val sample : 'a Distribution.t → ('a → 'b next) → 'b next
  val factor : float → (unit → 'b next) → 'b next
  val infer : ('a, 'b) model → 'a → 'b Distribution.t
end = struct ... end
```

Inference algorithm

- Run a set of n independent executions
- **sample**: draw a sample from a distribution
- **factor**: associate a score to the current execution
- Gather output values and score to approximate the posterior distribution

Importance sampling



Importance sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = ...

  let sample d k prob = assert false
  let factor s k prob = assert false

  let infer ?(n = 1000) m data = assert false
end
```

Importance sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = { id : int; particles : 'a particle array }
  and 'a particle = { k : 'a next; value : 'a option; score : float }
  ...

  let sample d k prob =
    let v = Distribution.draw d in
    k v prob

  let factor s k prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with score = s +. particle.score };
    k () prob
  ...
end
```

Importance sampling

infer.ml

```
module Importance_sampling = struct
  type 'a prob = { id : int; particles : 'a particle array }
  and 'a particle = { k : 'a next; value : 'a option; score : float }
  ...

  let exit v prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with value = Some v };
    prob

  let infer ?(n = 1000) model data =
    let particles = Array.make n { value = None; score = 0.; k = (model data) exit } in
    Array.iteri (fun i particle → ignore (particle.k { id = i; particles })) particles;
    let values = Array.map (fun p → Option.get p.value) particles in
    let logits = Array.map (fun p → p.score) particles in
    Distribution.support ~values ~logits
end
```

Coin

coin.ml

```
open Infer.Importance_sampling
```

```
let coin x =  
  let* z = sample (uniform ~a:0. ~b:1.) in  
  let* () = Cps_list.iter (observe (bernoulli ~p:z)) x in  
  return z
```

```
let _ =  
  let dist = infer coin [1; 1; 0; 0; 0; 0; 0; 0; 0; 0] in  
  let m, s = Distribution.stats dist in  
  Format.printf "Coin bias, mean:%f, std:%f@." m s
```

```
> dune exec ./examples/coin.exe
```

```
Coin bias, mean:0.247876, std:0.118921  
Beta(2+1, 8+1), mean:0.250000, std:0.120096
```

Particle filter (CPS)

BYO-PPL

Particle filter

basic.ml

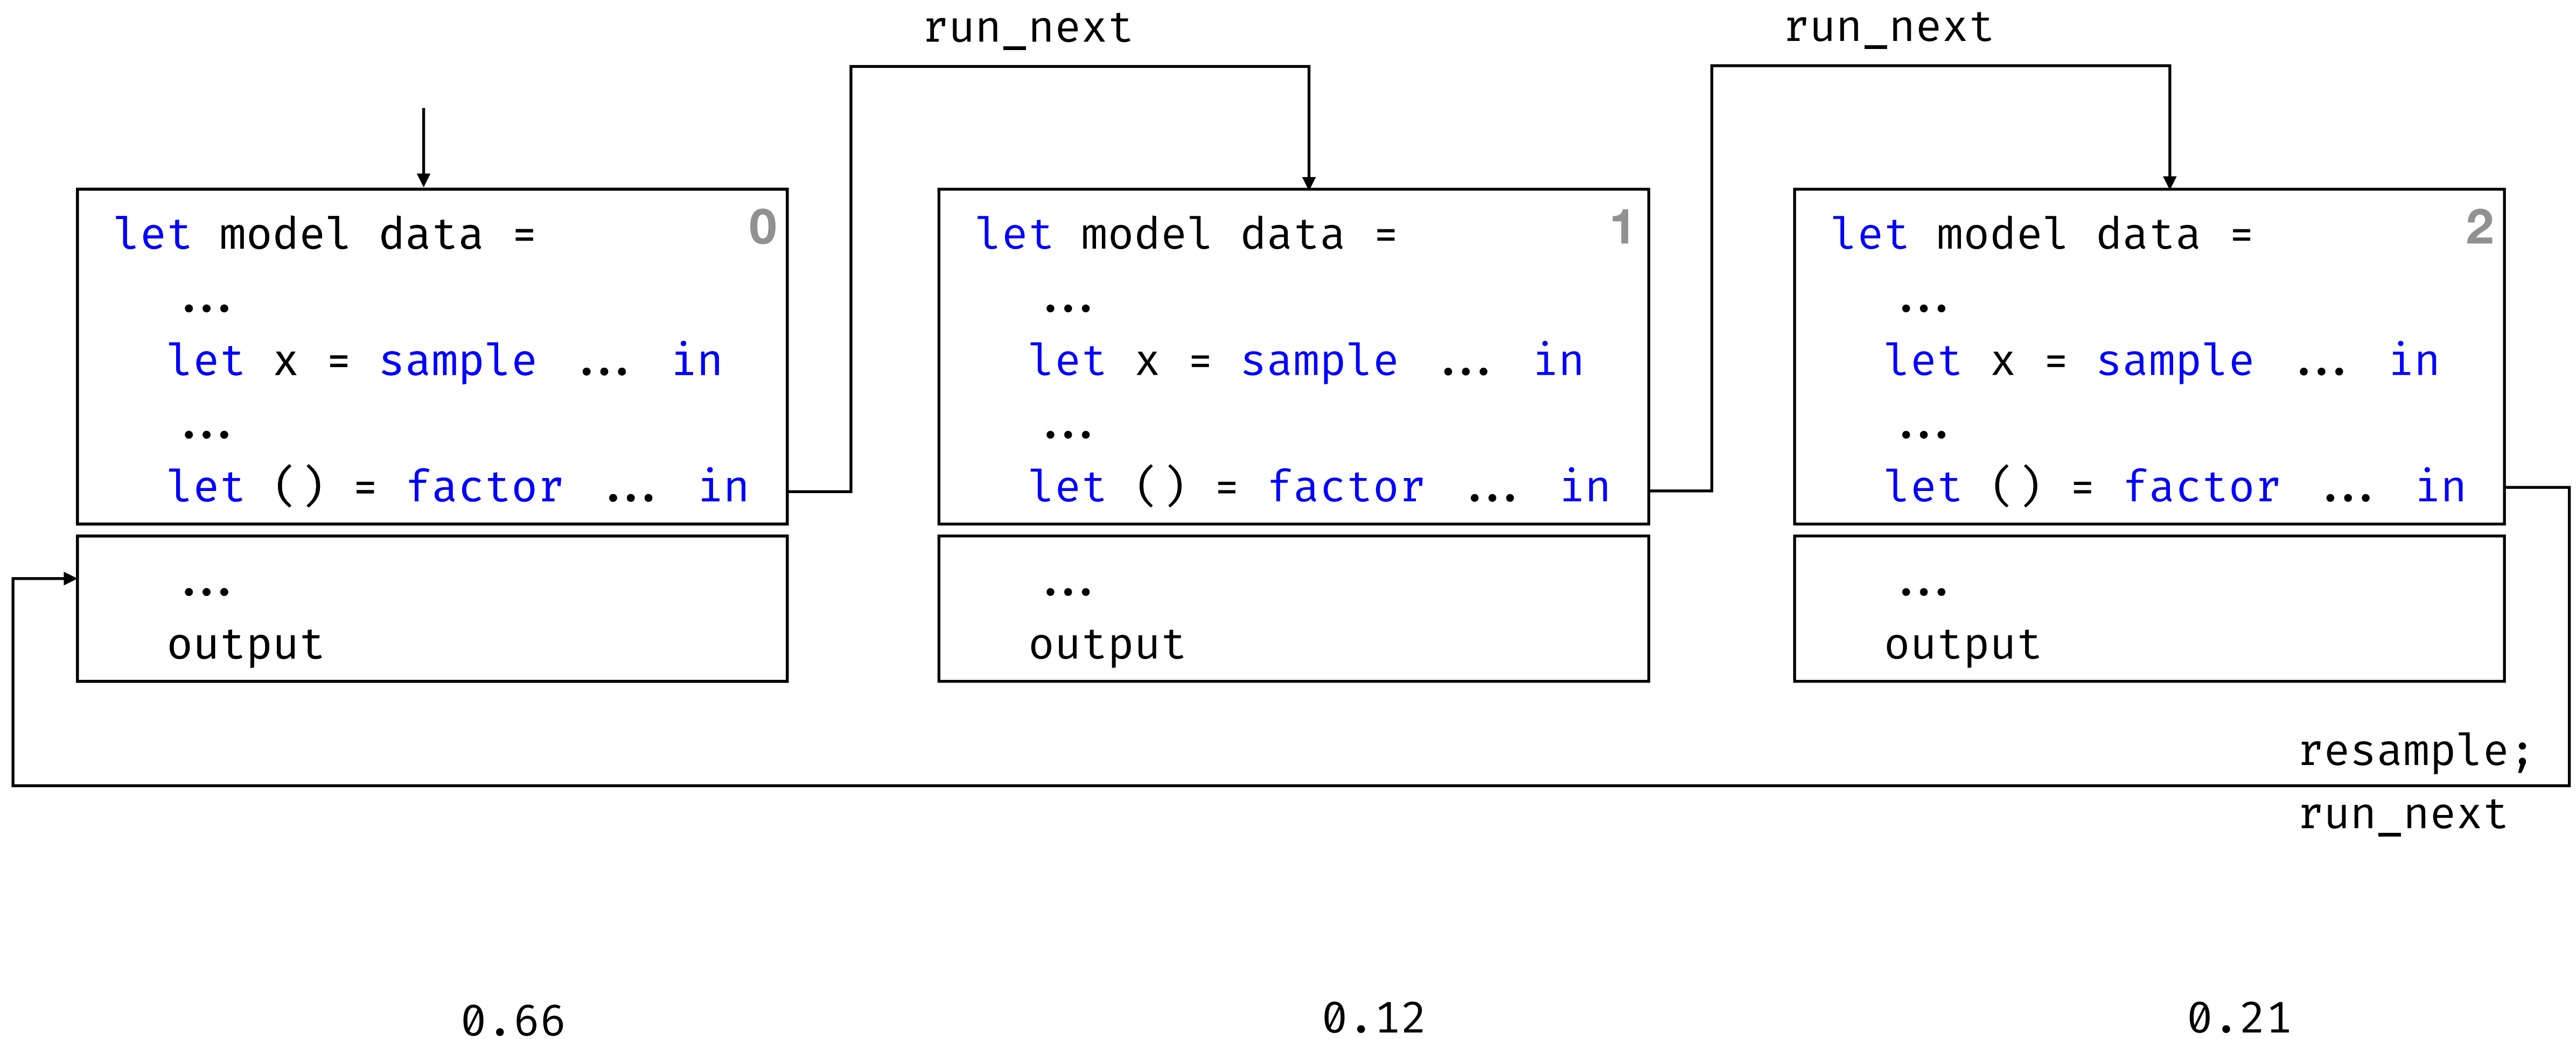
```
module Particle_filter = struct
  include Importance_sampling

  let resample particles = assert false
  let factor s k prob = assert false
end
```

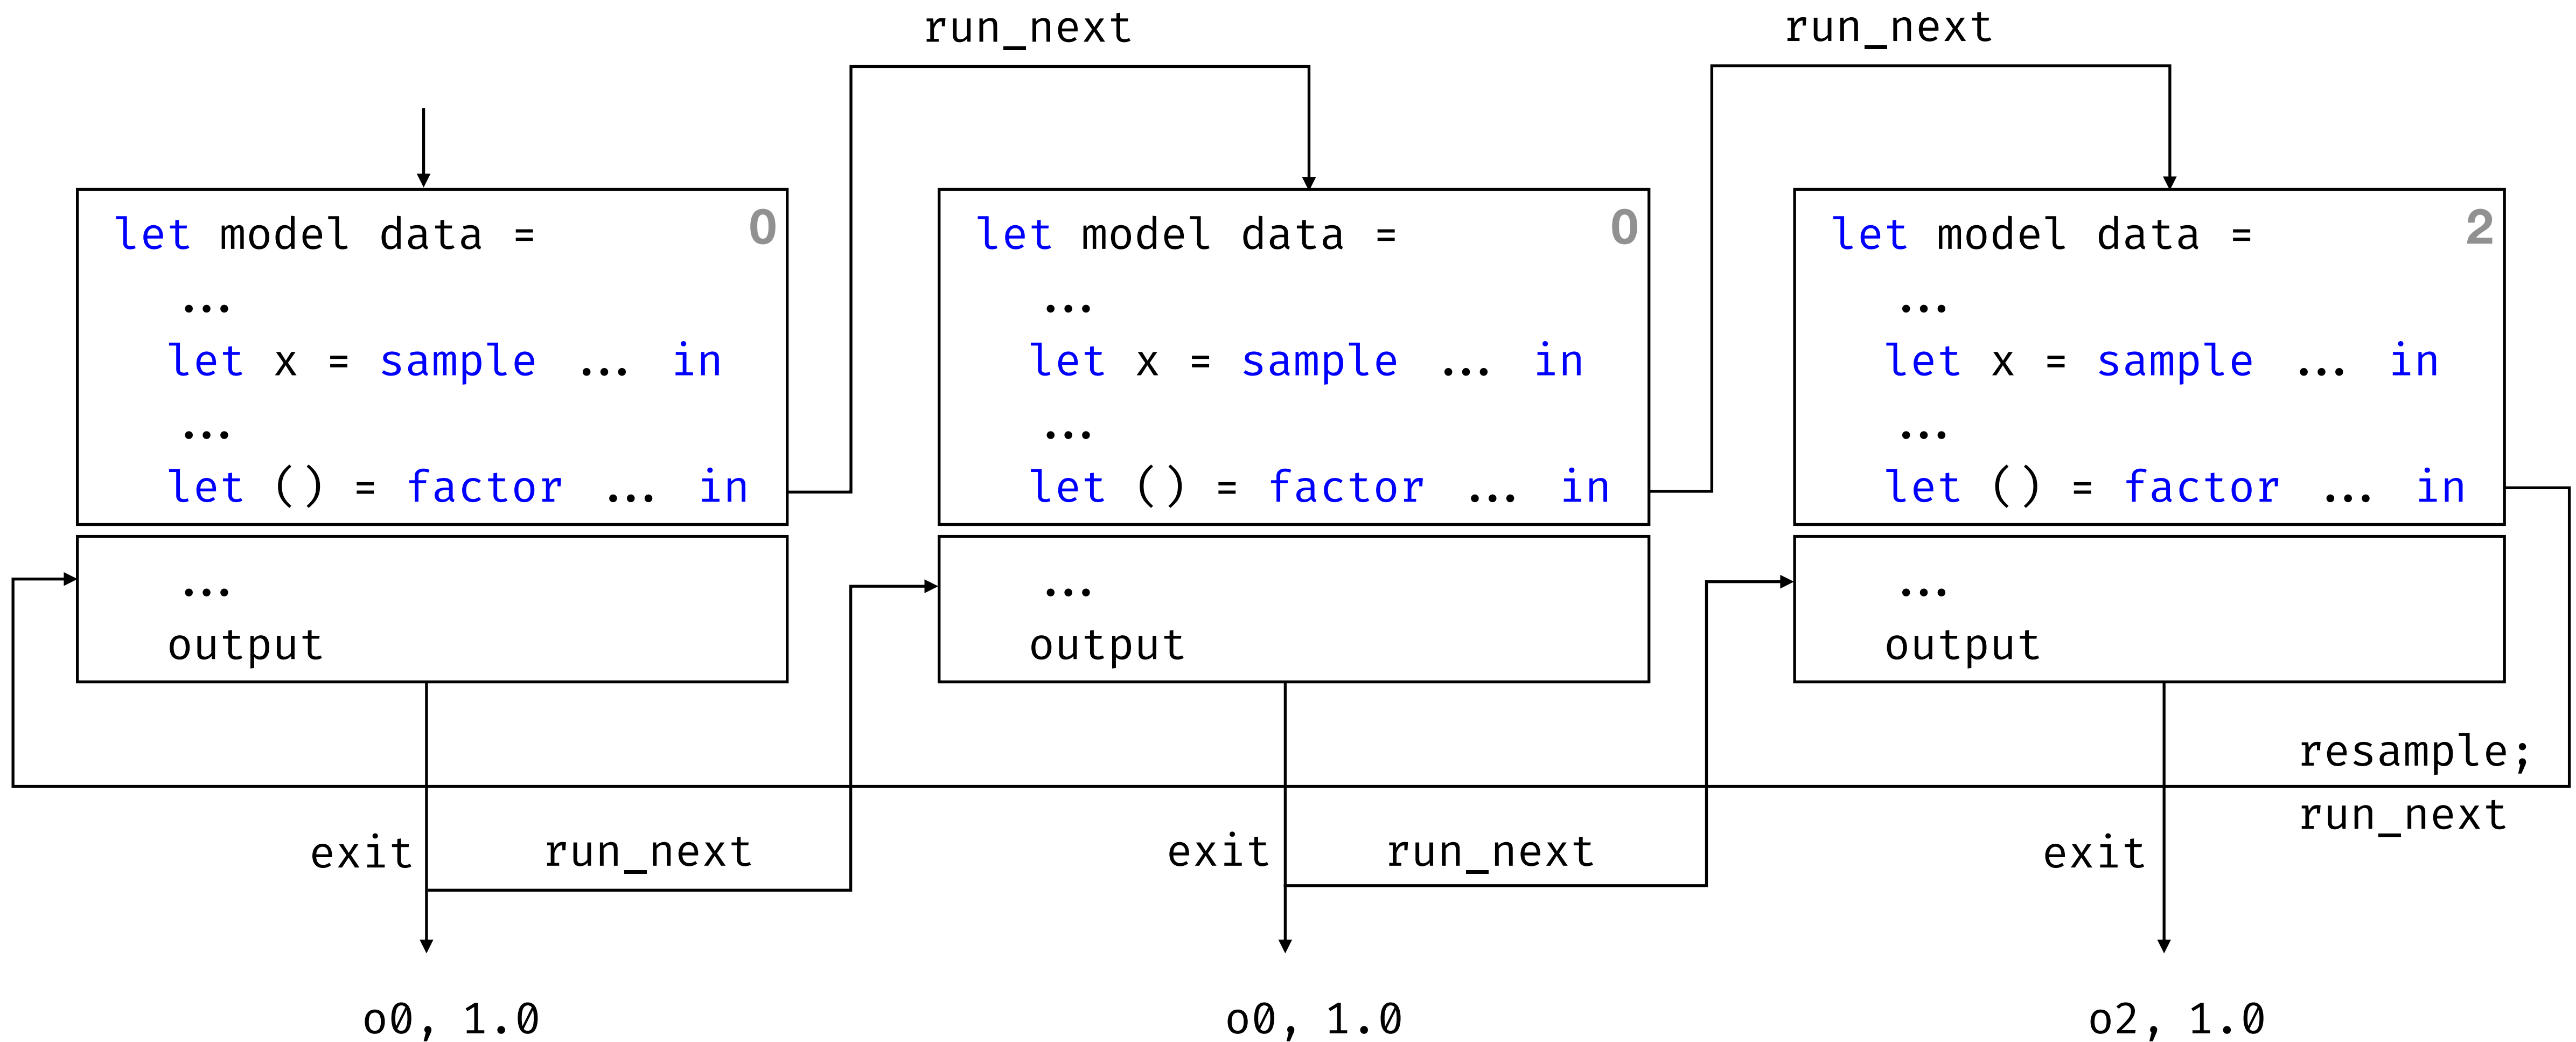
Inference algorithm : importance sampling, but...

- Add a resampling step at each **factor**
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution

Particle filter



Particle filter



Particle filter

infer.ml

```
module Particle_filter = struct
  include Importance_sampling

  let resample particles =
    let logits = Array.map (fun x → x.score) particles in
    let values = Array.map (fun x → { x with score = 0. }) particles in
    let dist = Distribution.support ~values ~logits in
    Array.iteri (fun i _ → particles.(i) ← Distribution.draw dist) particles

  let factor s k prob =
    let particle = prob.particles.(prob.id) in
    prob.particles.(prob.id) ← { particle with k = k (); score = s +. particle.score };
    prob
end
```

Particle filter

infer.ml

```
module Particle_filter = struct
  include Importance_sampling
  ...

  let infer ?(n = 1000) model data =
    let particles = Array.make n { value = None; score = 0.; k = (model data) exit } in
    while Array.for_all (fun p → Option.is_none p.value) particles do
      Array.iteri (fun i particle → ignore (particle.k { id = i; particles })) particles;
      resample particles
    done;
    let values = Array.map (fun p → Option.get p.value) particles in
    let logits = Array.map (fun p → p.score) particles in
    Distribution.support ~values ~logits
end
```

HMM: Hidden Markov Model

hmm.ml

```
open Infer.Particle_filter

let hmm prob data =
  let rec gen states data =
    match (states, data) with
    | [], y :: data → gen [ y ] data
    | states, [] → return states
    | pre_x :: _, y :: data →
      let* x = sample prob (gaussian ~mu:pre_x ~sigma:1.0) in
      let* () = observe prob (gaussian ~mu:x ~sigma:1.0) y in
      gen (x :: states) data
  in
  gen [] data


let _ =
  let data = Owl.Arr.linspace 0. 20. 20 ▷ Owl.Arr.to_array ▷ Array.to_list in
  let dist = Distribution.split_list (infer ~n:100 hmm data) in
  let m_x = List.map Distribution.mean dist in
  List.iter2 (Format.printf "%f >> %f") data m_x
```

HMM: Hidden Markov Model

```
› dune exec ./hmm.exe
```

```
0.000000 >> 0.000000
1.052632 >> 0.997546
2.105263 >> 2.300316
3.157895 >> 3.289649
4.210526 >> 4.857555
5.263158 >> 4.907179
6.315789 >> 6.254198
7.368421 >> 7.208341
8.421053 >> 8.432642
9.473684 >> 8.938143
10.526316 >> 9.555007
11.578947 >> 11.098199
12.631579 >> 12.823460
13.684211 >> 13.701444
14.736842 >> 14.934314
15.789474 >> 16.115058
```

...



Exercises

foo.ml

What if the particles do not terminate at the same time?

```
let foo () =  
  let* c = sample (bernoulli ~p:0.5) in  
  if c = 1 then  
    let* () = factor 1. in return c  
  else  
    return c
```

```
let _ =  
  let dist = infer foo () in  
  let { values; probs; _ } = get_support ~shrink:true dist in  
  Array.iteri (fun i x → Format.printf "%d %f@" x probs.(i)) values
```

Exercises

foo.ml

What if the particles do not terminate at the same time?

```
let foo () =  
  let* c = sample (bernoulli ~p:0.5) in  
  let* () = if c = 1 then factor 10. else skip in  
  return c  
  
let _ =  
  let dist = infer foo () in  
  let { values; probs; _ } = get_support ~shrink:true dist in  
  Array.iteri (fun i x → Format.printf "%d %f@." x probs.(i)) values
```

```
dune exec ./examples/foo.exe
```

```
Fatal error: exception Invalid_argument("option is None")
```


Exercise: Enumeration (CPS)

BYO-PPL

Enumeration

infer.ml

```
module : sig
  type 'a prob
  and 'a next = 'a prob → 'a prob
  and ('a, 'b) model = 'a → ('b → 'b next) → 'b next

  val sample : 'a Distribution.t → ('a → 'b next) → 'b next
  val factor : float → (unit → 'b next) → 'b next
  val infer : ('a, 'b) model → 'a → 'b Distribution.t
end = struct ... end
```

Inference algorithm: depth first exploration

- Maintain a continuation queue
- **sample**: pick one value (and score), push the other in the queue
- **factor**: update the score
- At the end of the model, store the pair (value, score), restart from the top of the queue

Inference comparison

BYO-PPL

Coin


coin.ml

```
let _ =  
  let data = List.init 1000 (fun i → i mod 2) in  
  let dist = infer coin data in  
  plot dist
```


Rejection Sampling

Very (very) slow!

Importance Sampling



Particle Filter



Particle impoverishment

HMM


hmm.ml

```
let _ =  
  let data = Owl.Arr.linspace 0. 100. 100 ▷ Owl.Arr.to_array ▷ Array.to_list in  
  let dist = Distribution.split_list (infer hmm data) in  
  plot (error (List.rev dist) data)
```


Rejection Sampling

Never terminate!

Importance Sampling



Particle Filter



Inference limitations

Rejection Sampling : Termination

- Only works if valid sample can be generated from the priors
- In practice: simple models with few discrete observations
- E.g., simple discrete models.

Importance Sampling : Weight collapse

- Score strictly decrease at each observation: eventually collapse
- In practice: continuous model to estimate fixed parameters from observations
- E.g., coin, linear regression

Particle Filter: Particle impoverishment

- Duplicate particle without resampling fixed parameter
- In practice: high-dimensional continuous models without fixed parameters
- E.g., hmm, tracker

Inference formalization

BYO-PPL

Sampler

Probabilistic semantics $G \vdash^P e : t$

- Expressions are interpreted as weighted samplers
- Draw random samples and track the execution weight
- Given an environment γ , $\llbracket e \rrbracket_\gamma = v, w$
- $\llbracket e \rrbracket : \Gamma \rightarrow V \times [0, \infty)$

$$\llbracket c \rrbracket_\gamma = c, 1$$

$$\llbracket x \rrbracket_\gamma = \gamma(x), 1$$

$$\llbracket \text{sample}(e) \rrbracket_\gamma = \text{draw}(\llbracket e \rrbracket_\gamma), 1$$

$$\llbracket \text{factor}(e) \rrbracket_\gamma = (), \llbracket e \rrbracket_\gamma$$

$$\begin{aligned} \llbracket \text{let } x = e_1 \text{ in } e_2 \rrbracket &= \text{let } v_1, w_1 = \llbracket e_1 \rrbracket_\gamma \text{ in} \\ &\quad \text{let } v_2, w_2 = \llbracket e_2 \rrbracket_{\gamma + [x \leftarrow v_1]} \text{ in} \\ &\quad v_2, w_1 \times w_2 \end{aligned}$$

Combine weights

Importance sampling

Inference algorithm

- Run a set of N independent executions
- `sample`: draw a sample from a distribution
- `factor`: associate a score to the current execution
- Gather output values and score to approximate the posterior distribution

$$\begin{aligned} \llbracket \text{infer}(e) \rrbracket_{\gamma} &= \text{let } [v_i, w_i = \{e\}_{\gamma}]_{1 \leq i \leq N} \text{ in} \\ &\quad \text{let } W = \sum_{i=1}^N w_i \text{ in} \\ &\quad \lambda U. \sum_{i=1}^N \frac{w_i}{W} \times \delta_{v_i}(U) \end{aligned}$$

Particle filter

Inference algorithm : importance sampling, but...

- Add a resampling step at each **factor**
- Compute the score of the particles to compute a distribution
- Re-sample a new set of particles from this distribution


Resampling

$$\begin{aligned} \{e\} &\propto \{\text{let } d = \text{infer}(e) \text{ in sample}(d)\} \\ \{\text{let } x = e_1 \text{ in } e_2\} &\propto \{\text{let } d = \text{infer}(e_1) \text{ in let } x = \text{sample}(d) \text{ in } e_2\} \\ \llbracket \text{infer}(\text{let } x = e_1 \text{ in } e_2) \rrbracket &= \llbracket \text{infer}(\text{let } d = \text{infer}(e_1) \text{ in let } x = \text{sample}(d) \text{ in } e_2) \rrbracket \end{aligned}$$

Checkpoints

BYO-PPL

Big-step semantics with checkpoints



Reduction rules: stop evaluation when r is true

Big-step semantics with checkpoints

$$\frac{e \rightarrow^{false} true, 1 \quad e_1 \rightarrow^r e'_1, w_1}{\text{if } e \text{ then } e_1 \text{ else } e_2 \rightarrow^r e'_1, w_1}$$

$$\frac{e \rightarrow^{false} false, 1 \quad e_2 \rightarrow^r e'_2, w_2}{\text{if } e \text{ then } e_1 \text{ else } e_2 \rightarrow^r e'_2, w_2}$$

$$\frac{e_1 \rightarrow^{true} e'_1, w_1}{\text{let } x = e_1 \text{ in } e_2 \rightarrow^{true} \text{let } x = e'_1 \text{ in } e_2, w_1}$$

$$\frac{e_1 \rightarrow^{false} v_1, w_1 \quad e_2 \rightarrow^r e'_2, w_2}{\text{let } x = e_1 \text{ in } e_2 \rightarrow^r e'_2, w_1 \times w_2}$$

Combine weights

$$\frac{e \rightarrow^{false} \mu, 1}{\text{sample}(e) \rightarrow^{true} \text{draw}(\mu), 1}$$

$$\frac{e \rightarrow^{false} s, 1}{\text{factor}(e) \rightarrow^{true} (), s}$$

Stop!

Particle filter

$$\frac{[e_i \rightarrow^{r_i} e'_i, w_i]_{1 \leq i \leq N} \quad \bigvee_{1 \leq i \leq N} r_i \quad \rho = \text{Cat} \left(\{e'_i, w_i\}_{1 \leq i \leq N} \right) \quad [\text{draw}(\rho)]_{1 \leq i \leq N} \Rightarrow \mu}{[e_i]_{1 \leq i \leq N} \Rightarrow \mu}$$

Resampling

$$\frac{[e_i \rightarrow^{false} v_i, w_i]_{1 \leq i \leq N} \quad W = \sum_{i=1}^N w_i}{[e_i]_{1 \leq i \leq N} \Rightarrow \lambda U. \sum_{i=1}^N \frac{w_i}{W} \delta_{v_i}}$$

Gather the results

$$\frac{[e]_{1 \leq i \leq N} \Rightarrow \mu}{e \Rightarrow^N \mu}$$

Launch N particles

References

WebPPL

Noah Goodman and Andreas Stuhlmüller

<http://webppl.org/>

The Design and Implementation of Probabilistic Programming Languages

Noah Goodman and Andreas Stuhlmüller

<http://dippl.org/>

An Introduction to Probabilistic Programming

Jan-Willem van de Meent, Brooks Paige, Hongseok Yang, Frank Wood

<https://arxiv.org/abs/1809.10756>

Embedded probabilistic domain-specific language HANSEI

Oleg Kiselyov, Chung-chieh Shan

<https://okmij.org/ftp/kakuritu/Hansei.html>

BYO-PPL

Build Your Own Probabilistic Language

- Clone the repo: `git clone https://github.com/mpri-probprog/byo-ppl-22-23.git`
- Install the dependencies: `opam install . --deps-only`
- Build the project: `dune build`
- Test an example: `dune exec ./examples/funny_bernoulli.exe`

Implemented as an OCaml embedded domain specific language (eDSL)

- `Distribution`: small library of probability distributions and basic statistical functions.
- `Basic`: basic inference algorithms (rejection sampling inference sampling)
- `Infer`: inference algorithms for models written in Continuation Passing Style (CPS).
- `Cps_operators`: syntactic sugar to write CPS style probabilistic models.
- `Utils`: missing utilities functions used in other modules.